Beauty Contest

June 18, 2014

Abstract

Pithy, concise and informative. May bring the reader to tears due to the beauty of it.

1 Introduction

2 Methods

The availability of large data sets for building regression models to predict the bacterial counts in beach water is both an opportunity and a challenge.

2.1 Data Sources

Possibly move this to the end of the section

Which sites

Where are they

What specific sources sources of data (plug EnDDAT)

Will include a map and tables

2.2 Listing of specific statistical techniques

Fourteen different regression modeling techniques were considered. Each technique uses one of five modeling algorithms: GBM, the adaptive lasso, the genetic algorithm, PLS, or sparse PLS. Each technique is aplied to either continuous or binary regression and to either modeling, or variable selection only.

Continuous or binary regression

The goal of predicting exceednaces of the water quality standard was approached in two ways: one was to predict the bacterial concentration and then compre the prediction to a threshold. The other was to predict the state of a binary indicator, which is coded as zero when the concetration is below the standard and

one when the concentration exceeds the standard. Techniques taking the former approach are continuous regression techniques, those taking the latter approach are binary regression techniques.

Weighting of observations in binary regression

A weighting scheme was implemented for some of the binary regression techniques. In the weighting scheme, observations were given weights equal to the number of standard deviations the observed concentration was from the regulatory threshold of 235 CFU/100 mL. Any technique that was implemented with this weighting scheme was separately implemented without any weighting of the observations. The techniques are then labeled weighted and unweighted, respectively.

Modeling or selection only

Some methods are labeled "select", which means that they are used for variable selection only. In these cases, once the variables are selected, the regression model is estimated using ordinary least squares for the continuous regression techniques, or ordinary logistic regression for the binary regression techniques.

2.2.1 GBM

GBM refers to the gradient boosting machine (GBM) of **?**. A GBM model is a so-called random forest model - a collection of many regression trees. Prediction is done by averaging the outputs of the trees. Two GBM-based techniques are explored - we refer to them as GBM and GBMCV. The difference is in how the optimal number of trees is determined - GBMCV selects the number of trees in a model using leave-one-out CV, while GBM uses the so-called out-of-bag (OOB) error estimate. The CV method is much slower (it has to construct as many random forests as there are observations, while the OOB method only requires computing a single random forest) but GBMCV should more accurately estimate the prediction error. All the GBM and GBMCV models share the following settings:

Number of trees: 10000

Shrinkage parameter: 0.0005

Minimum observations per node: 5

Depth of each tree: 5

Bagging fraction: 0.5

2.2.2 Adaptive Lasso

The adaptive lasso **?** is a regression method that simultaneously selects relevant predictors and estimates their coefficients by adding a penalty to the sum of the squared residuals. For linear regression, the adaptive

lasso estimates $\hat{\boldsymbol{\beta}}$ minimize the criterion $\sum_{i=1}^{n}(y_i - X_i\beta)^2 + \lambda\sum_{j=1}^{p}\frac{|\beta_j|}{|\tilde{\beta}_j|^\gamma}$, where λ is a tuning parameter and $\tilde{\boldsymbol{\beta}}$ is a consistent estimate of the regression coefficients.

In this work, γ is set to one, $\tilde{\boldsymbol{\beta}}$ are estimated individually by a univariate linear or logistic regression (it is necessary to estimate the coefficients individually because there are usually more covariates than observations) and the adaptive lasso tuning parameter λ is selected to minimize the AICc [**?**].

Three of the modeling techniques were based on the adaptive lasso - one fo

2.2.3 Genetic algorithm

The genetic algorithm [**?**] is a variable-selection method that works by analogy to natural selection, where so-called chromosomes represent regression models. A variable is included in the model if the corresponding element of the chromosome is one, but not otherwise. Chromosomes are produced in successive generations, where the first generation is produced randomly and subsequent generations are produced by combining chromosomes from the current generation, with additional random drift. The chance that a chromosome in the current generation will produce offspring in the next generation is an increasing function of its fitness. The fitness of each chromosome is calculated by the corrected Akaike Information Criterion (AICc) **??**.

The implementations in this study used 100 generations, with each generation consisting of 200 chromosomes. The genetic algorithm method GALM is the default for linear regression modeling in Virtual Beach [**?**]. The study also investigates two genetic algorithm methods for logistic regression: one weighted (GALogistic-weighted) and one unweighted (GALogistic-unweighted).

2.2.4 PLS

Partial least squares (PLS) regression is a tool for building regression models with many covariates [**?**]. PLS works by decomposing the covariates into mutually orthogonal components, with the components then used as the variables in a regression model. This is similar to principal components regression (PCR), but the way PLS components are chosen ensures that they are aligned with the model output. On the other hand, PCR is sometimes criticised for decomposing the covariates into components that are unrelated to the model's output.

To use PLS, one must decide how many components to use in the model. The technique used in this study follows the method described in **?**, using the PRESS statistic to select the number of components.

2.2.5 SPLS

Sparse PLS (SPLS) is introduced in **?**. The SPLS method combines the orthogonal decompositions of PLS with the sparsity of lasso-type variable selection. To do so, SPLS uses two tuning parameters: one that controls the number of orthogonal components and one that controls the lasso-type penalty. The optimal

parameters are those that minimize the mean squared prediction error (MSEP) over a two-dimensional grid search. The MSEP is calculated by 10-fold cross-validation. Two techniques utilizing SPLS were

2.3 Implementation for beach regression

The response variable for our continuous regression models is the natural logarithm of the E. coli concentration. For the binary regression models, the response variable is

Include a table with pre/post processing discussion

This includes tuning of parameters

Some specific data issues because we are estimating a threshold exceedence

2.4 Cross Validation

Assessment of the modeling techniques is based on their performance in predicting exceedances of the regulatory standard. Two types of cross validation was used to measure the performance in prediction: leave-one-out (LOO) and leave-one-year-out (LOYO). In LOO CV, one observation is held out for validation while the rest of the data (the model data) is used to train a model. The model is used to predict the concentration of that held out observation, and the process is repeated for each observation. Each cycle of LOYO CV holds out one year's worth of data for validation instead of a single observation. It is intended to approximate the performance of the modeling technique under a typical use case: a new model is estimated before the start of each annual beach season and then used for predicting exceedances during the season. That year's data is then added to the dataset to estimate a model for the next beach season. The LOYO models in this study were estimated using all the available data, even that from future years - so for instance the 2012 models were estimated using the 2010-2011 and 2013 data

Some methods also used cross-validation internally to select tuning parameters. In those cases the internal CV was done by partitioning the model data, leaving out one partition at a time. This process is separate from - and does not affect - the CV to assess predictive performance.

2.5 Performance Metrics

How did we evaluate the performance of each technique on all the different data sets

- for all cases —> AUC (ROC curve)
- continuous variables using PRESS (skill –> Like Nash-Sutcliffe/R^2 over the fitted data)
- True/False Positives/Negatives (needs a threshold)
- Which variables are selected for models where variable reduction takes place

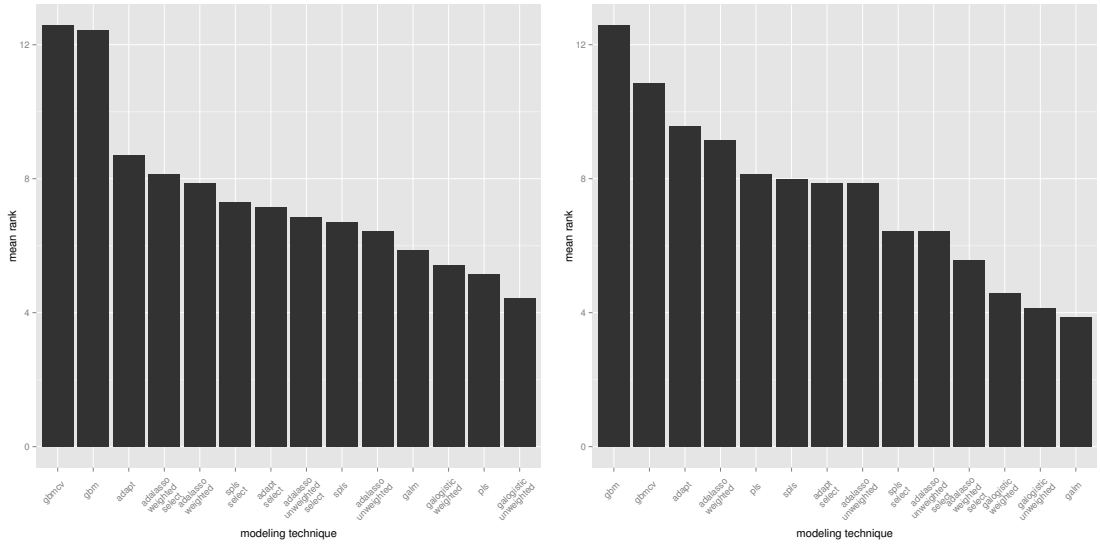

Figure 3.1: Mean ranks of the modeling techniques across the seven sites. At left are the mean ranks under leave-one-out cross validation, at the right are the mean ranks from leave-one-year-out cross validation.

- challenge regarding the fact that different variables are selected in each fold. Maybe use frequencies?
- also the number of variables selected (metric of complexity)

OPTIONAL

- AIC/BIC? –> not the same model in each fold so maybe not possible
- Maybe some form of confusion matrices – perhaps a grid of them with or without companion variance plots or other estimates of the range of results

3 Results

The area under the ROC (AUROC) curve assesses how accurately the predictions of the left-out observations are sorted. The methods are ranked at each site by AUROC and a mean rank (across sites) is computed for each method. The mean ranks are plotted in

- Performance over prediction from cross validation
- Maybe anecdotal showing fit over data set

4 Discussion

In general, the GBM, and GBMCV, and AL techniques produced comparable results that were superior to the other techniques in terms of predictive performance. Since the GBMCV models take much longer to

compute than the others, we will not include them in our more detailed analysis of the modeling results.

Which type of model is generally the best?

Under what conditions do some outperform others?

Relative value of overall best model versus methods that help trim variables? e.g. how valuable is it to reduce number of predictors? Further, which variables get cut? Expensive ones? Cheap ones?

How important is computational expense? Only an issue for model fitting — not prediction, but worth quantifying. E.g. if GBM with cross validation takes hours, how much better?

Model tuning for GBM versus GBM-CV –> notes on how GBM is faster with similar performance (e.g. CV is overkill maybe)

5 Acknowledgments

6 References

References